

Tim Miller
15.071 | HW #3

Problem 1a

Running a linear regression yields a model R^2 of 0.8937. See sample output below (not including everything because it is very long). Some coefficients are reported as NA because they overlap with other variables in the dataset so do not contribute to the model. For example TotalBsmtSF is just the sum of the other Basement Sq Ft metrics (finished / unfinished) so having TotalBsmtSF does not add anything.

```
> pia.mod = lm(data=train, SalePrice ~ .)
> summary(pia.mod)

Call:
lm(formula = SalePrice ~ ., data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-279704  -11312   -474    10735   120142

Coefficients: (7 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -3.512e+05  9.047e+05  -0.388  0.697928
MSZoningRL      8.152e+03  3.894e+03   2.093  0.036451 *
MSZoningRM      6.757e+03  4.501e+03   1.501  0.133475
LotFrontage    -1.737e+02  3.848e+01  -4.514  6.78e-06 ***
LotArea         6.255e-02  1.089e-01   0.575  0.565698
StreetPave     -1.904e+03  1.263e+04  -0.151  0.880195
AlleyNo Alley   2.470e+03  3.528e+03   0.700  0.483909
AlleyPave       5.447e+02  5.565e+03   0.098  0.922029
LotShapeMod+ IR  9.075e+03  3.474e+03   2.612  0.009066 **
LotShapeReg     1.139e+03  1.413e+03   0.806  0.420203
LandContourHLS  2.038e+04  4.383e+03   4.649  3.58e-06 ***
LandContourLow  1.383e+04  5.502e+03   2.514  0.012034 *
LandContourLv1  1.370e+04  3.105e+03   4.413  1.08e-05 ***
LotConfigCu1DSac 2.787e+03  2.974e+03   0.937  0.348715
LotConfigFR2    -7.953e+03  3.684e+03  -2.159  0.031006 *
LotConfigFR3    -8.545e+03  8.400e+03  -1.017  0.309170
LotConfigInside -9.686e+02  1.590e+03  -0.609  0.542598
LandSlopeNot Gtl 6.337e+03  3.580e+03   1.770  0.076834 .
...
...
...
...
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 23850 on 1794 degrees of freedom
Multiple R-squared:  0.8937,    Adjusted R-squared:  0.8822
F-statistic: 77.74 on 194 and 1794 DF,  p-value: < 2.2e-16
```

Problem 1b

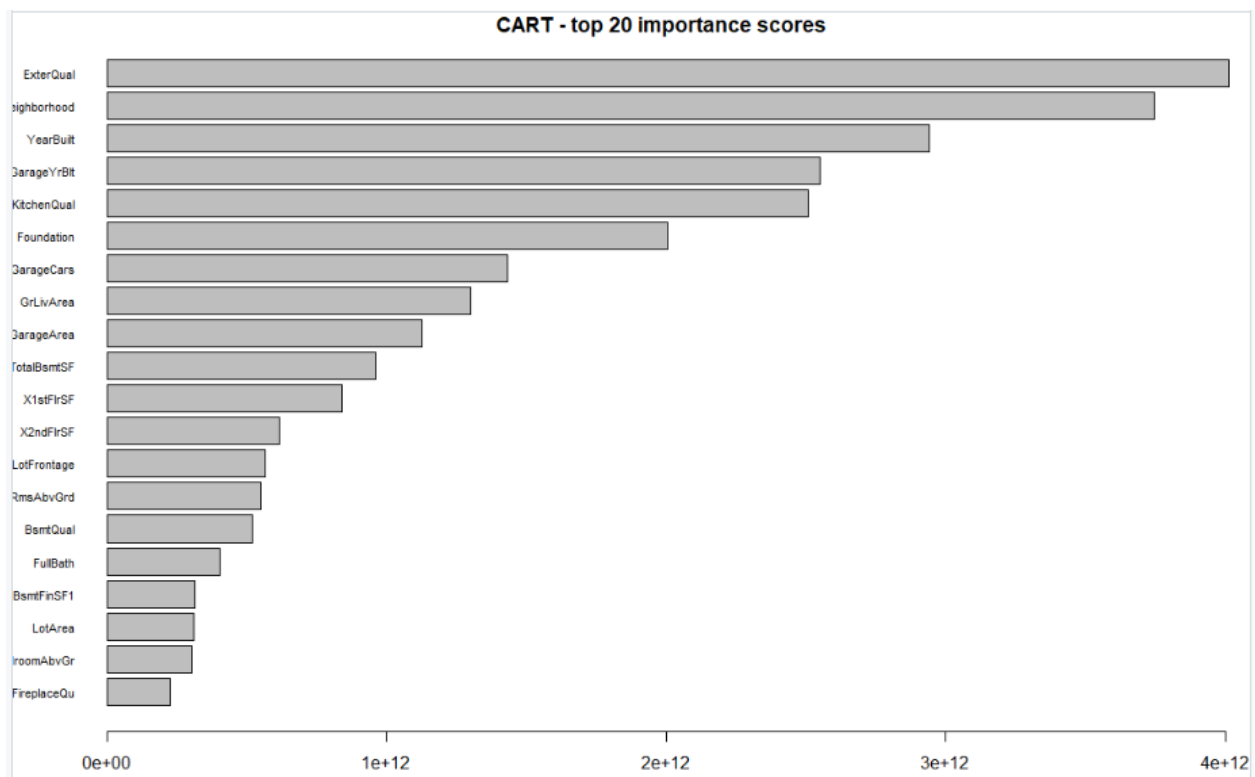
We want to find the cp value with the min RMSE. When we run the below code, we find the **cp value is 0.000296**.

```
> cv.results[which.min(cv.results$RMSE),]
      cp      RMSE Rsquared      MAE  RMSESD RsquaredSD  MAESD
287 0.000296 30913.59 0.8036813 21541.93 2275.525 0.02855522 1360.901
```

Code to train the model with above cp value:


```
#train model
cart.mod = rpart(SalePrice ~ ., data = train, cp=0.000296)
prp(cart.mod)
```

Below we see the top 20 most important variables. Clear trend is that variables related to the quality and durability of the home are shown as important predictors for sale price (e.g., External Material Quality, Year Built, Foundation, Kitchen Quality). Perhaps Ames is known for some weather events and people want a sturdy house. Additionally, people probably want to buy a house that they don't need to fix up. I would imagine lower quality houses go for lower prices because the buyer and owner know money needs to go into the house to make improvements.



Problem 1c

Plotting RMSE vs mtry we find that that value of mtry is 28.

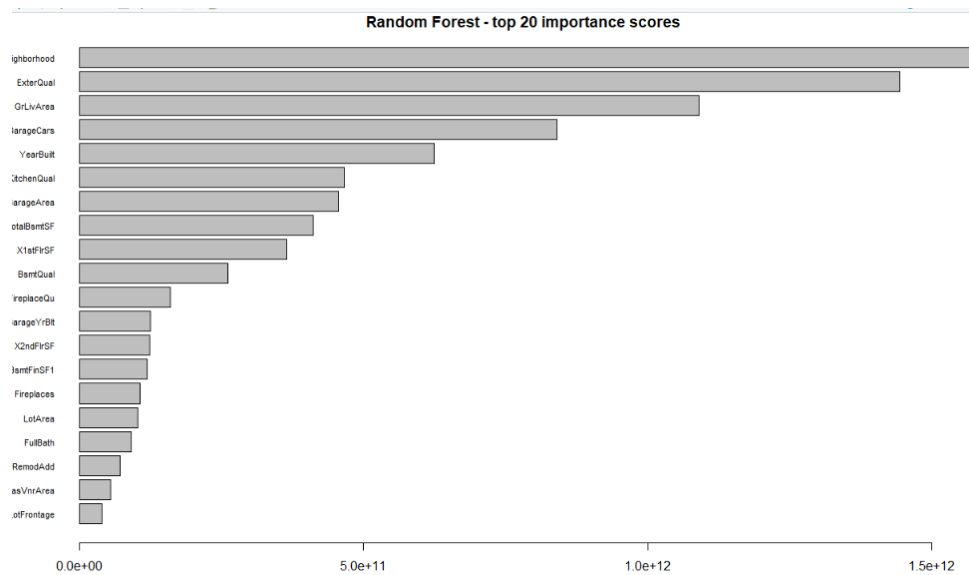


Problem 1d

Using the below code to extract the final model:

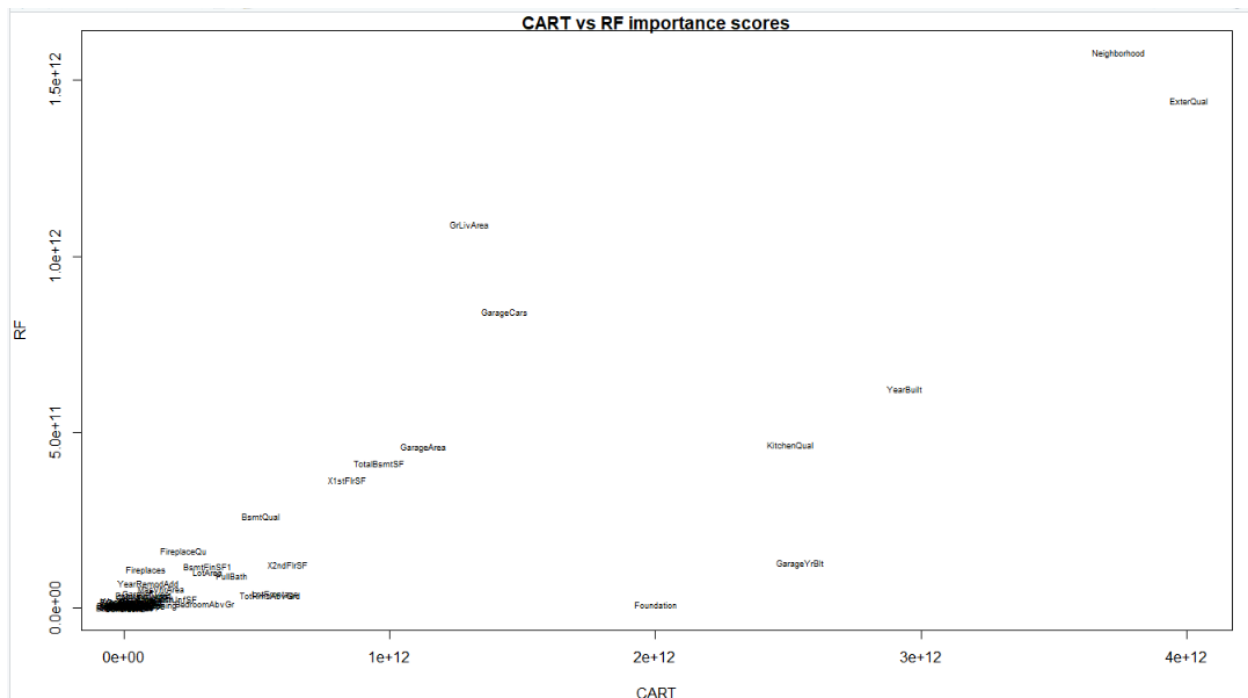
```
#extract best model|  
mod.rf = train.rf.oob$finalModel
```

We get the chart below for top 20 importance variables. The variables that appear in the Random Forest model seem well placed and intuitive. Most variables corresponding to the size and quality of the house.



Problem 1e

Interesting to see GrLivArea higher up in this model vs. the CART model. Intuitively I would have guessed that GrLivArea would be high up given people pay more money for bigger houses (i.e., surprised it was lower down in the CART model). Furthermore, foundation does not crack the top 20 for the Random Forest, whereas it is #4 important variable for CART. But that does not surprise me too much. I wonder how many perspective buyers *actually* look at the home foundation as they are making the bid on the home. Perhaps this is a random correlation picked up in the CART model that is corrected for in the Random Forest model.



Problem 1f

As expected the R^2 goes up from Linear Regression, to CART, to Random Forest. Also we can confirm that the OSR^2 for CART drops, but does not for Random Forest. A potential explanation here is that the CART model overfit to certain variables e.g., Foundation because they were more prevalent in the train set, but less in the test set. So when we run the model on the test set the predictive power is not as good because of the overfitting.

Linear regression

- In sample $R^2 = 0.8937$
- $OSR^2 = 0.8854796$

CART

- In sample $R^2 = 0.9108464$
- $OSR^2 = 0.7977558$

```
9 # Make predictions on test and train sets
0 PredictTrain = predict(cart.mod, newdata = train)
1 PredictTest = predict(cart.mod, newdata = test)
2
3 # Calculate R-Squared and OSR-Squared
4 SSTTrain = sum((train$SalePrice - mean(train$SalePrice))^2)
5 SSETrain = sum((PredictTrain - train$SalePrice)^2)
6 R2_CART <- 1 - SSETrain/SSTTrain
7 SSTTest = sum((test$SalePrice - mean(test$SalePrice))^2)
8 SSETest = sum((PredictTest - test$SalePrice)^2)
9 OSR2_CART <- 1 - SSETest/SSTTest
0
1 R2_CART
2 OSR2_CART
3
```

Random Forest

- In sample $R^2 = 0.9559223$
- $OSR^2 = 0.8851491$

```
#RANDOM FOREST

# Make predictions on test and train sets
PredictTrain = predict(mod.rf, newdata = train)
PredictTest = predict(mod.rf, newdata = test)

# Calculate R-Squared and OSR-Squared
SSTTrain = sum((train$SalePrice - mean(train$SalePrice))^2)
SSETrain = sum((PredictTrain - train$SalePrice)^2)
R2_RF <- 1 - SSETrain/SSTTrain
SSTTest = sum((test$SalePrice - mean(test$SalePrice))^2)
SSETest = sum((PredictTest - test$SalePrice)^2)
OSR2_RF <- 1 - SSETest/SSTTest

R2_RF
OSR2_RF
```

Problem 1g

Problem 1h

- ### Problem 2a

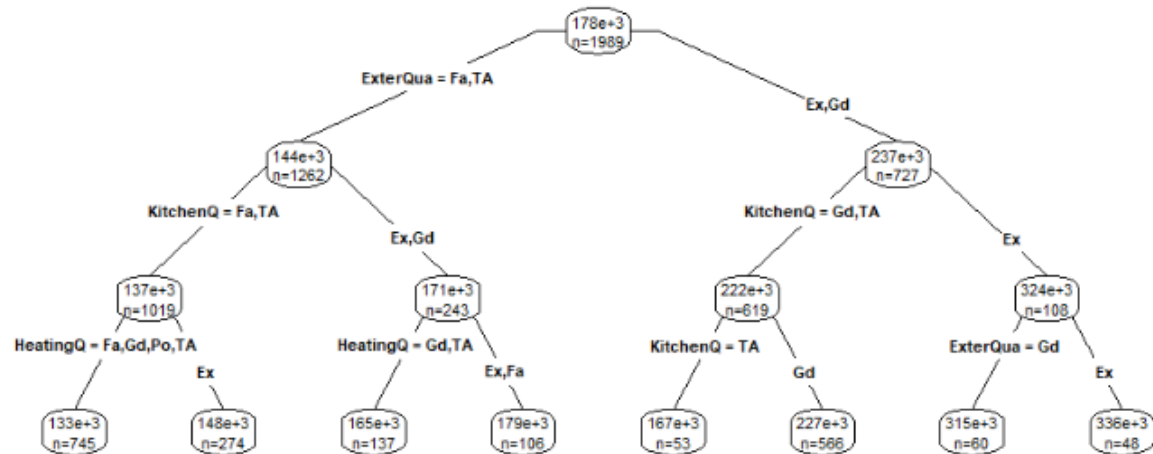
For example, unordered sorting just orders alphabetically

Whereas ordered sorting sorts by the order we provided

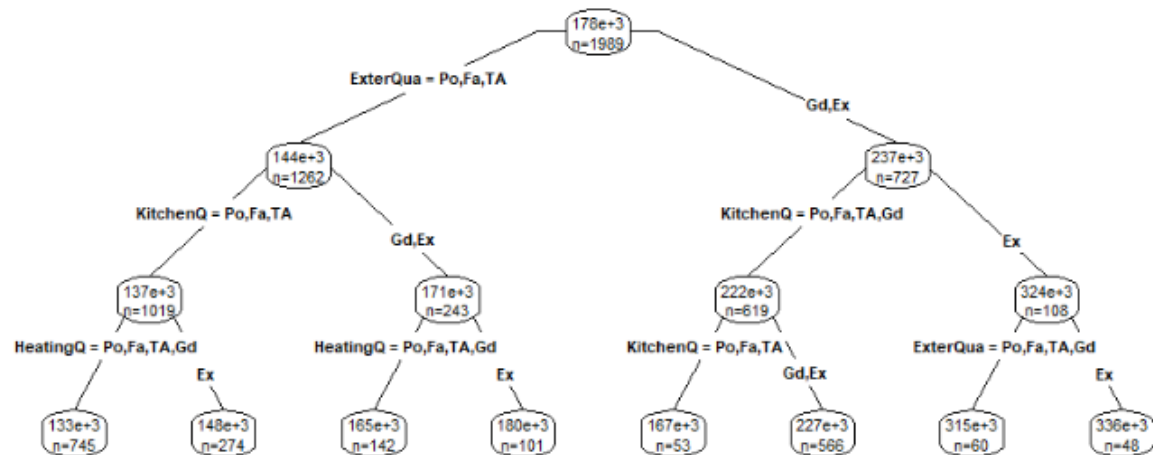
Problem 2b

Furthermore, we see that the Po quality tag shows up much more often for the ordered tree than the unordered tree. This is evident in the first split where we see Po show up for the ordered tree and not the unordered tree.

A: Unordered factors



B: Ordered factors



Problem 2c

The ordered factor tree makes much more sense. As stated in 2b, we want to group the quality tags that are most similar together. Furthermore, we want to make sure that the Po quality tag is not lost in the analysis because it is at the end of the alphabet.

I would choose to use the ordered tree model in my analysis.

Problem 2d

Based on R^2 it looks like TreeA fits the data better as it has a slightly higher R^2 .

Tree A $R^2 = 0.5641749$


```
# Make predictions on test and train sets
PredictTrain = predict(treeA, newdata = trainA)

# Calculate R-Squared and OSR-Squared
SSTTrain = sum((trainA$SalePrice - mean(trainA$SalePrice))^2)
SSETrain = sum((PredictTrain - trainA$SalePrice)^2)
R2_RF_TREEA <- 1 - SSETrain/SSTTrain
R2_RF_TREEA
```

(Untitled) :

Tree B $R^2 = 0.5641458$

```
4 |
5 # Make predictions on test and train sets
6 PredictTrain = predict(treeB, newdata = trainB)
7
8 # Calculate R-Squared and OSR-Squared
9 SSTTrain = sum((trainB$SalePrice - mean(trainB$SalePrice))^2)
10 SSETrain = sum((PredictTrain - trainB$SalePrice)^2)
11 R2_RF_TREEB <- 1 - SSETrain/SSTTrain
12 R2_RF_TREEB
13
```

(Untitled) :

Problem 2e

It looks like Tree B performs better on the test set as it has a slightly higher OSR².

Tree A OSR² = 0.5337154

```
# Calculate R-Squared and OSR-Squared
PredictTest = predict(treeA, newdata = testA)
SSTTest = sum((testA$SalePrice - mean(testA$SalePrice))^2)
SSETest = sum((PredictTest - testA$SalePrice)^2)
OSR2_RF <- 1 - SSETest/SSTTest

OSR2_RFA
```

Tree B OSR² = 0.5358374

```
PredictTest = predict(treeB, newdata = testB)
SSTTest = sum((testB$SalePrice - mean(testB$SalePrice))^2)
SSETest = sum((PredictTest - testB$SalePrice)^2)
OSR2_RF <- 1 - SSETest/SSTTest

OSR2_RFB
```

Problem 2f

Ordered factoring reduces the run time for CART models. By introducing logical ordering, we reduce the search time of the algorithm to find the best split. For example, with the models we developed in problem 2. With the ordered values, we help group the similar quality levels close to each other. Intuitively we know that Po and Fa should be grouped close together and we get a better model *when* they are grouped together. By doing this for the computer, it does not have to go out and try additional combinations to discover that through trial and error.

Problem 2g

An example that might work is gradient of hair color as a predictor of standardized test scores.

Hair color gradient has clear and logical ordering: you would want to order from dark to light (or vice versa).

But, intuitively, I don't see a scenario where hair color matters at all for predicting standardized test scores. Therefore we would violate the second condition.